

# FORMALIZATION OF DISTRIBUTED SYSTEMS WITH SEMANTIC INTEROPERABILITY

Leon BRÂNZAN\*  
[leon.branzan@doctorat.utm.md](mailto:leon.branzan@doctorat.utm.md)

March, 29th, 2024

## Abstract

Distributed software systems have been designed, studied, and implemented for decades, yet problems with their development, deployment, and maintenance persist even today. Attempts at formalizing the crucial concepts of distributed systems often lead nowhere or fail outright, as is demonstrated in this article.

A hypothesis is then proposed: using mathematical models dealing with semantics of interoperability of systems it is possible to develop a better understanding of distributed computing using not the objects within the system, but the relations between these objects. The article describes a use case for applying semantic analysis to solve persisting problems with industrial systems.

Viable solutions to these problems are then suggested, borrowed from well-formalized mathematical theories, such as domain theory and category theory. The article attempts to partially answer the questions it poses using “semantic interoperability”—the property of a notation to have different formal definitions of the same concept be fully interchangeable in the context of a unifying formal description.

**Keywords:** *denotational semantics; category theory; distributed systems; network architecture.*

## Contents

<b>Introduction</b>	<b>2</b>
<b>Centralized networks in a decentralized architecture</b>	<b>2</b>
<b>Computation and storage</b>	<b>3</b>
<b>Semantic interoperability</b>	<b>4</b>
<b>Categories</b>	<b>5</b>
<b>Conclusions</b>	<b>6</b>

---

\*Departamentul Informatică și Ingineria Sistemelor, Facultatea Calculatoare, Informatică și Microelectronică, Universitatea Tehnică a Moldovei, Chișinău, Republica Moldova

## Introduction

On March 5th, 2024, Facebook, Instagram, Threads, and several other communication services developed by Meta suffered a global outage, resulting in millions of people losing access to their accounts for the duration of the outage [1]. It is speculated that the core issue lied with the company's internal network infrastructure. This was not the first time such an event took place. On October 4th, 2021, a massive outage disrupted Meta's services globally. In a post-mortem [2] Meta's engineers narrowed down the issues to an error with DNS configuration. Such episodes are not specific to Meta, nor are they rare. Distributed systems "are hard" [3].

The problem here is not with the systems themselves. Engineers dedicate a lot of man-hours making sure such systems stay available and reliable. Rather, the problem is the over-reliance on private systems for global communication, oftentimes in critical moments [4]. The Internet was initially developed as an open network, and communication was and still is its primary function. Communication channels should not be gated by private entities. Specifically, international communication should be subject to international law. It is difficult to police private companies outside of one's authority in cases when said companies fail to comply with local laws, especially since it is much easier for private entities to deny their services rather than litigate. Global communication relies almost entirely on independent entities, critical infrastructure that millions depend on is out of most people's reach.

When large companies like Amazon, Meta etc. cannot fix the issues with stability of their systems, it is a good indicator that these problems cannot be fixed just with money. Networks created on top of the Internet must be aware of the underlying infrastructure and replicate its most important properties.

## Centralized networks in a decentralized architecture

The World Wide Web was initially devised as a network for exchanging texts between scientific institutions [5]. But even then, its designer Tim Berners-Lee considered it just an initial step, and the future of the network to be in machines communicating with other machines [6]. XML was one of the formats intended to bridge that gap. But instead, social networks took over and are still omnipresent. Even though the phenomenon, which is today known under the name "Web 2.0", came about as the result of multiple efforts to democratize the process of publishing content on web sites [7], its arguably biggest impact lies in concentrating most human communication within several private centralized networks; the other important aspect of this process is the noticeable effect that mainstream advertising practices have on the nature of content published on the biggest social networks [8].

As illustrated in "The Semantic Web" [5], the future of the World Wide Web looked different to its creators and early adopters. That future was based on evolving the way various entities within the Web exchanged information (a good example of that view is Szabo's seminal work on secure information exchange on public networks [9]), closer in its architecture to a huge peer-to-peer network. Today, most of the Web's communication is orchestrated by huge centralized systems. All the while peer-to-peer networks are reserved for hobbyists; federated communication networks outside email are considered niche.

One would expect a significant shift in this paradigm with the advent of the Internet of Things. This model seems to mimic the kind of architecture described in "The Semantic Web." It is often just a replication of already existing architectures (most of which originate in social networks; for example, Software-as-a-Service monetization schemes being prevalent in IoT solutions, with companies denying services at their convenience [10, 11, 12]). On the surface, IoT should be an integral part of the Internet, seemingly inseparable. If Internet service is available, the device should be fully functional. Currently, when the company ceases to service or update their product, it becomes unusable.

There are attempts at changing the status quo. Several protocols and implementations have been proposed in the last 5–10 years, with their adoption lagging behind. ActivityPub is one such example [13], its mission being rebuilding the Web as a decentralized system, as it was imagined before "everything got locked down into a handful of walled gardens". Notable implementations are Mastodon and Blue Sky. Another example, which aims to "radically change the way Web applications work today, resulting in true data ownership as well as improved privacy", is Solid—a project headed by Berners-Lee himself [14].

An alternative approach to solving this problem would be to change the underlying software in a way that would compel users to change their behavior. This can be traced in how such changes in behavior

were happening before: a new way to interact with the system would be identified, adopted by a small group of people, gain critical mass and then explode in popularity<sup>1</sup>. Web 2.0 is a good illustration of this process. From an engineering perspective developing a new system for enacting a similar shift would be an insurmountable task even for a group of developers, let alone one person. A different strategy needs to be adopted. Instead of creating an unproven system and then expecting it to eventually do what it was designed to do, another approach would be to first prove that such a system is:

- a. viable;
- b. capable of changes expected of it;
- c. will work according to its specification.

A rigorous proof would require a precise formal description (existing or new). Considering the nature of distributed systems, describing a proof concerning relations between objects in such a system would require either developing a new general formal language or finding a way of unifying existing formalisms (syntactic models) for describing distributed systems (e. g. Erlang’s “Actor” model [15]). One possible approach could be based on generalizing different models using a more abstract notion of relations between objects.

## Computation and storage

Distributed systems research revolves around computation and storage. A lot of mainstream research is focused on large text processing, consensus protocols, data replication, and performance<sup>2</sup>. In the curriculum the gap between early research papers and modern research topics is noticeably large (this could be attributed, at least in part, to COVID-19). Still, a lot of effort has been spent over the decades to formalize certain aspects of distributed systems, sometimes with unexpected results.

One rather contentious topic in distributed systems design is the CAP<sup>3</sup> theorem [16]. The original conjecture [17] states, in plain terms, that it is possible to “have at most two of these properties for any shared-data system:

- consistency;
- availability;
- tolerance to network partitions.”

The theorem garnered a lot of exposure and critique [18, 19]. Famously, Kleppmann once noted that, “the CAP theorem is too simplistic and too widely misunderstood to be of much use for characterizing systems. Therefore I ask that we retire all references to the CAP theorem, stop talking about the CAP theorem, and put the poor thing to rest” [20]. Invariant Confluence [21] is proposed by many as a viable alternative.

Both approaches formalize data consistency when synchronized over a network. The other important topic is network consensus. Two major results in this domain are Paxos [22] and Viewstamped Replication [23]. While the former is more widely known, there are well-documented attempts at making it “more approachable” for developers, garnering the algorithm a reputation of being difficult to implement. The more recent Raft algorithm implements Viewstamped Replication with several new features [24] and should be considered over older consensus algorithms.

The research discussed above shows that distributed systems formalization is an important topic, and that there are many problems in distributed systems design, that are still unsolved. One of the preliminary conclusions of this paper is the assumptions that the Internet’s properties and processes need to be formalized to make their replication and adoption in higher-level architectures more widespread and systematic. The focus, then, is on how distributed systems orchestrate communication between its actors, depending on the architecture (decentralized, federated, peer-to-peer, or hybrid). The intuition is: to see how that could work it would be beneficial to look at distributed systems from “a bird’s eye view”—one level of abstraction higher.

<sup>1</sup>This process is sometimes referred to as “Crossing the Chasm”, a phrase coined by G. A. Moore.

<sup>2</sup>This assumption is based on MIT’s 2023 curriculum for the “Distributed Systems” course, which is available at [pdos.csail.mit.edu](https://pdos.csail.mit.edu).

<sup>3</sup>CAP is an acronym that stands for “consistency, availability, partitions”.

## Semantic interoperability

The “semantics” of a system is its behavior. From a broad point of view, semantics and realization are aspects of the same situation: semantics is the problem of system analysis; while realization is the problem of system synthesis [25]. In general, semantics are separated into three major classes:

1. **Operational.** Meanings for program phrases defined in terms of the steps of computation they can take during program execution.
2. **Axiomatic.** Meanings for program phrases defined indirectly via the axioms and rules of some logic of program properties.
3. **Denotational.** Concerned with giving mathematical models of programming languages. Meanings for program phrases defined abstractly as elements of some suitable mathematical structure.

When attempting to look at concrete things closely, having a formal way of abstracting all the details would help immensely. Researchers often turn to formalization while looking for viable solutions to concrete problems. When formalizing distributed informational systems one important property that needs to be preserved is semantic interoperability—“what is sent is what is understood” (for the purposes of this text semantic interoperability is defined as in [26]).

Consider the issues that could arise from fragmenting distributed networks along the connections inside them. Such division will inevitably make a network heterogeneous, which immediately leads to several problems that need to be addressed [27]: data incompatibility, the need for APIs at each point of connection, new metadata schemas etc.

One concrete example would be the Internet of Things. While each device can be connected to the internet and have a well-defined human-machine interface, things can quickly break down when attempting to make two devices communicate with each other (see “smart objects” [28] for one proposed solution).

Another example is programming language interoperability. Languages have the extra burden of syntactic interoperability and semantic interoperability. Existing solutions often revolve around creating a language extension or a language framework to overcome this issue. Other solutions attempt to formalize the higher-level concepts of a language (see “linear language interoperability” [29] for one proposed solution).

One example of an interoperable formal description of a property of a distributed system is MixT [30]—a C++-derived transaction language for concurrent computations, that enables its type system (and the compiler by extension) to catch incorrect formalisms. It is partially based on the concept of full abstraction borrowed from denotational semantics.

Listing 1 represents a language embedding designed to solve issues with concurrent mutation—the iterator in the fragment on the left can be invalidated by one thread while another is accessing its value. To combat this, a “transaction block” is introduced with the `mixt_method` declaration. Such blocks are context-aware, which allows them to safely merge concurrent operations, ensuring causal consistency.

(b) Code embedded within a MixT structure

(a) Unsafe C++ code

```
var iterator = users,
while (iterator.isValid()) {
    log.append(iterator->v
    .inbox.insert(post)),
    iterator = iterator->next
}
```

```
class User {
    Handle<set<string>>, causal, supports<insert>> inbox;
};

class group {
    RemoteList<Handle<user, causal>, linearizable> users;
    Handle<Log, eventual, supports<append>> log;
    mixt_method(add_post) (post) mixt_captures(users, log)
    (...);
};
```

Listing 1: MixT language embedding[30]

To summarize, denotational semantics serve “to specify programming language constructs in as abstract and implementation-independent way as possible: in this way one may gain insight into the fundamental concepts underlying programming languages, their inter-relationships, and (sometimes) new ways of realising those concepts in language designs [31].”

## Categories

Another mathematical theory that deals in abstractions and uses mathematical concepts to describe all sorts of entities is category theory. It is being widely used in modern research to explain different phenomena that were hard or impossible to describe using other methods.

It is outside of the scope of this paper to include a section on the basics of category theory, for a detailed illustrated introduction refer to [32]. Below is a diagram describing three main mathematical properties of an object that qualify it as a category (see fig. 1).




Figure 1: Illustration of three basic properties of a category [32].

Even though category theory is often called “abstract nonsense,” it has real application in many domains. One example is using the Yoneda lemma—one of the foundational theorems of category theory. To understand the core of the Yoneda lemma: in simple terms, imagine a deck of playing cards. If one person picks a random card and asks another person to guess it by asking questions about it, it would be possible to pin down the card in a finite number of questions (for example, questions like “is it a spade?”, “is it higher than a 10?” etc. will eventually lead to the correct card by elimination). In even plainer terms, it is possible to define an object by its relation to other objects definitively. Figure 2 presents a visual reference for the “Inverted spectrum” problem [33].



Figure 2: Illustration of the “inverted spectrum” problem [33].

If two people look at the same set of objects and are asked to name the objects’ colors, both will name the same colors, even when one of the persons has color vision deficiency. That person will have grown up with knowing a certain color as “red,” even though it would not necessarily qualify as “red” on the color spectrum. The Yoneda lemma gives a solution to this. Figure 3 shows the illustrated solution to the “Inverted spectrum” problem.

By transforming the color spectrum into a distorted space, the problem can be reframed as a mathematical problem. Any point of that space can be determined in terms of its relations with all other points of the space. Any point considered being in the red spectrum will have only one way of defining it in terms of all other points. Thus, any person not seeing it in the red part of the spectrum can be identified.

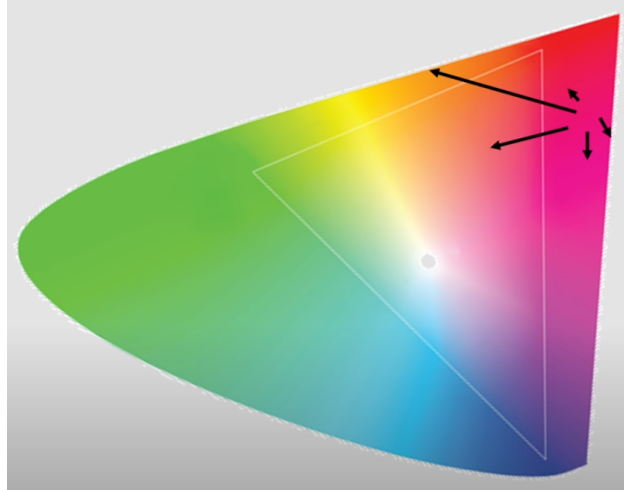


Figure 3: Human color perception as a distorted space [34].

## Conclusions

Another important concept of category theory is that of a functor. Without going into too much detail, a functor represents transformations between categories in the same way that functions are transformations between objects within a category (in this case—a set). Functors have one important property that is pertinent to this paper’s subject. Applying a functor to a composition of two transformations is equivalent to applying a functor to each transformation separately and then composing the results [35]. This is illustrated as a diagram below (see fig. 4).




Figure 4: Functor preserving commuting and composition [35].

This can be applied to describing how two heterogeneous systems interface with each other. If there are two systems that are connected to each other somehow, having some operations available in one system, it should give the same result in both cases:

- \* Operations are combined (read, performed one immediately after another “piping” the intermediate result between them), and then the end result is transported to the other system.
- \* Operations themselves are transported into the other system and combined there.

Meaning, that if both systems formally guarantee that the operations and objects possess certain properties, functors have the explanatory power to guarantee consistency of data between heterogeneous systems.

This is just one example of how category theory could be leveraged for designing and describing complex interconnected systems. In that the author of this paper agrees with Joseph Goguen, “computing science is



very fragmented, with many different sub-disciplines having many different schools within them. Hence, we badly need the kind of conceptual unification that category theory can provide [36].

Of course, the path to unification need not lie in category theory necessarily. But it would be a good first step, because concepts developed using category theory are easily generalized and can rely on many mathematical proofs to ensure that what they describe they do with enough precision. To quote Scott and Strachey again, “there are many different languages adequate for conveying the same concepts (e.g., binary, octal, or decimal numerals). Even in the same language many different expressions can denote the same concepts (e.g.,  $2 + 2$ ,  $4$ ,  $1 + 1 + 1 + 1$ , etc.). The problem of explaining these equivalences of expressions (whether the same or different languages) is one of the tasks of semantics and is much too important to be left to syntax alone. Besides, the mathematical concepts are required for the proof that the various equivalences have been correctly described [31].”

## References

- [1] A. Rita, “Facebook Service Outage Postmortem,” Medium. Accessed: Apr. 10, 2024. [Online]. Available: [medium.com](https://medium.com).
- [2] S. Janardhan, “More details about the October 4 outage,” Engineering at Meta. Accessed: Apr. 10, 2024. [Online]. Available: [engineering.fb.com](https://engineering.fb.com).
- [3] K. Kingsbury, “The network is reliable,” Aphyr. Accessed: Apr. 10, 2024. [Online]. Available: [aphyr.com](https://aphyr.com).
- [4] D. Davies, “Social media’s role in Jan. 6 was left out of the final report,” NPR. Accessed: Apr. 10, 2024. [Online]. Available: [npr.org](https://npr.org).
- [5] T. Berners-Lee, “Information Management: A Proposal,” CERN, TBL-900620, May 1990.
- [6] T. Berners-Lee, J. Hendler, and O. Lassila, “The Semantic Web: A New Form of Web Content that is Meaningful to Computers will Unleash a Revolution of New Possibilities,” in *Linking the World’s Information*, 1st ed., O. Seneviratne and J. Hendler, Eds., New York, NY, USA: ACM, 2001, pp. 91–103. doi: 10.1145/3591366.3591376.
- [7] M. Sykora, “Web 1.0 to Web 2.0: an observational study and empirical evidence for the historical r(evolution) of the social web,” *IJWET*, vol. 12, no. 1, p. 70, 2017, doi: 10.1504/IJWET.2017.084024.
- [8] Y. Walter, “Artificial influencers and the dead internet theory,” *AI & Soc*, Feb. 2024, doi: 10.1007/s00146-023-01857-0.
- [9] N. Szabo, “Formalizing and Securing Relationships on Public Networks,” *First Monday*, vol. 2, no. 9, Sep. 1997, [Online]. Available: [firstmonday.org](https://firstmonday.org).
- [10] P. Acoba, “Tesla allegedly remotely unlocks Model 3 owner’s car, uses smart summon to help repo agent,” *Tire Meets Road*. Accessed: Apr. 10, 2024. [Online]. Available: [tiremeetsroad.com](https://tiremeetsroad.com).
- [11] J. M. Gitlin, “John Deere relents, says farmers can fix their own tractors after all,” *Ars Technica*. Accessed: Apr. 10, 2024. [Online]. Available: [arstechnica.com](https://arstechnica.com).
- [12] J. Hamzelou, “A brain implant changed her life. Then it was removed against her will,” *MIT Technology Review*. Accessed: Apr. 10, 2024. [Online]. Available: [technologyreview.com](https://technologyreview.com).
- [13] C. Lemmer-Webber, J. Tallon, E. Shepherd, A. Guy, and E. Prodromou, “ActivityPub,” W3C, Recommendation REC-activitypub-20180123, Jan. 2018. Accessed: Apr. 10, 2024. [Online]. Available: [w3.org](https://w3.org).
- [14] E. Mansour et al., “A Demonstration of the Solid Platform for Social Web Applications,” in *Proceedings of the 25th International Conference Companion on World Wide Web—WWW ’16 Companion*, Montreal, Quebec, Canada: ACM Press, 2016, pp. 223–226. doi: 10.1145/2872518.2890529.
- [15] C. Hewitt, “Actor Model of Computation.” *arXiv*, Nov. 07, 2010. doi: 10.48550/arXiv.1008.1459.

- [16] S. Gilbert and N. Lynch, “Brewer’s Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services,” SIGACT News, vol. 33, no. 2, pp. 51–59, Jun. 2002, doi: 10.1145/564585.564601.
- [17] E. A. Brewer, “Towards robust distributed systems (abstract),” in Proceedings of the nineteenth annual ACM symposium on Principles of distributed computing, in PODC ’00. Portland Oregon USA: ACM, Jul. 2000, p. 7. doi: 10.1145/343477.343502.
- [18] M. Kleppmann, “A Critique of the CAP Theorem.” arXiv, Sep. 18, 2015. Accessed: Mar. 09, 2024. [Online]. Available: [arxiv.org/abs/1509.05393](https://arxiv.org/abs/1509.05393).
- [19] D. Tornow, “The CAP Theorem. The Bad, the Bad, & the Ugly,” DTornow. Accessed: Apr. 04, 2024. [Online]. Available: [blog.dtornow.com](https://blog.dtornow.com).
- [20] M. Kleppmann, “Please stop calling databases CP or AP,” Martin Kleppmann. Accessed: Apr. 10, 2024. [Online]. Available: [martin.kleppmann.com](https://martin.kleppmann.com).
- [21] P. Bailis, A. Fekete, M. J. Franklin, A. Ghodsi, J. M. Hellerstein, and I. Stoica, “Coordination avoidance in database systems,” Proc. VLDB Endow., vol. 8, no. 3, pp. 185–196, Nov. 2014, doi: 10.14778/2735508.2735509.
- [22] L. Lamport, “The Part-Time Parliament,” ACM Transactions on Computer Systems, vol. 16, no. 2, pp. 133–169, May 1998, doi: 10.1145/279227.279229.
- [23] B. Oki and B. Liskov, “Viewstamped Replication: A New Primary Copy Method to Support Highly-Available Distributed Systems,” in Proceedings of the Seventh Annual ACM Symposium on Principles of Distributed Computing, in PODC ’88. Toronto, Ontario, Canada: Association for Computing Machinery, 1988, pp. 8–17. doi: 10.1145/62546.62549.
- [24] D. Ongaro and J. Ousterhout, “In Search of an Understandable Consensus Algorithm,” in 2014 USENIX Annual Technical Conference, Stanford University: USENIX Association, Jun. 2014, pp. 305–319. Accessed: Apr. 10, 2024. [Online]. Available: [usenix.org](https://usenix.org).
- [25] J. A. Goguen, “Semantics of Computation,” Computer Science & Engineering. Accessed: Apr. 10, 2024. [Online]. Available: [cseweb.ucsd.edu](https://cseweb.ucsd.edu).
- [26] E. Miscena and C. de Guerdavid, Eds., “European Interoperability Framework: Implementation Strategy.” National Interoperability Framework Observatory, Mar. 23, 2017. Accessed: Apr. 10, 2024. [Online]. Available: [joinup.ec.europa.eu](https://joinup.ec.europa.eu).
- [27] B. H. De Mello et al., “Semantic interoperability in health records standards: a systematic literature review,” Health Technol., vol. 12, no. 2, pp. 255–272, Mar. 2022, doi: 10.1007/s12553-022-00639-w.
- [28] O. Novo and M. D. Francesco, “Semantic Interoperability in the IoT: Extending the Web of Things Architecture,” ACM Trans. Internet Things, vol. 1, no. 1, pp. 1–25, Feb. 2020, doi: 10.1145/3375838.
- [29] G. Scherer, M. New, N. Rioux, and A. Ahmed, “FabULous Interoperability for ML and a Linear Language.” arXiv, Apr. 12, 2018. Accessed: Feb. 29, 2024. [Online]. Available: [arxiv.org/abs/1707.04984](https://arxiv.org/abs/1707.04984).
- [30] M. Milano and A. C. Myers, “MixT: a language for mixing consistency in geodistributed transactions,” in Proceedings of the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation, Philadelphia PA USA: ACM Press, Jun. 2018, pp. 226–241. doi: 10.1145/3192366.3192375.
- [31] D. Scott and C. Strachey, “Toward a Mathematical Semantics for Computer Languages,” Oxford University, Oxford University Computer Laboratory Programming Research Group, Monograph PRG-6, Oct. 1971. [Online]. Available: [cs.ox.ac.uk](https://cs.ox.ac.uk).
- [32] N. Tsuchiya and H. Saigo, “A relational approach to consciousness: categories of level and contents of consciousness,” Neuroscience of Consciousness, vol. 2021, no. 2, Aug. 2021, doi: 10.1093/nc/niab034.



- [33] A. Byrne, “Inverted Qualia,” Stanford Encyclopedia of Philosophy. Accessed: Apr. 10, 2024. [Online]. Available: [plato.stanford.edu](https://plato.stanford.edu).
- [34] A. Maier, “Category Theory for Neuroscience (pure math to combat scientific stagnation)”. Accessed: Apr. 10, 2024. [Presentation]. Available: [youtube.com](https://youtube.com).
- [35] B. Ahrens, C. Kapulkin, and M. Shulman, “Univalent categories and the Rezk completion,” Math. Struct. Comp. Sci., vol. 25, no. 5, pp. 1010–1039, Jun. 2015, doi: 10.1017/S0960129514000486.
- [36] J. A. Goguen, “A categorical manifesto,” Math. Struct. Comp. Sci., vol. 1, no. 1, pp. 49–67, Mar. 1991, doi: 10.1017/S0960129500000050.